

# ml\_dev - Machine Learning Development Image

---

This image supports the use of a Docker container for the development of Machine Learning projects in an Ubuntu environment.

## Table of Contents

1. Installed core components
  2. Creating a new Machine Learning development container
  3. Working with an existing Machine Learning development container
  4. Best practices
- 

## 1. Installed core components

With the following command you can check in detail which software versions are included in the Docker image:

```
apt list --installed
```

Version 1.1.0

Component	Version	Remark	Status
Anaconda			removed
asdf	v0.8.1-a1ef92a		new
cURL	7.77.0		new
Docker Compose	1.29.2		upgrade
Docker Engine	20.10.7		upgrade
dos2unix	7.4.2		new
Eclipse	2021-06-R		upgrade
GCC & G++	10.3.0		new
Git	2.32.0		upgrade
Gradle	7.1.1		upgrade
htop	3.0.5		new
ImageMagick	7.1.0-2		upgrade
Java	16.0.1	openjdk	upgrade
Julia	1.6.1		new

Component	Version	Remark	Status
Python	3.9.6		upgrade
- alpha_vantage	2.3.1		upgrade
- fbprophet			removed
- fire	0.4.0		upgrade
- Keras	2.4.3		
- matplotlib	3.4.2		upgrade
- numpy	1.19.5		upgrade
- pandas	1.3.0		upgrade
- pip	21.1.3		upgrade
- prophet	1.0.1		new
- pystan	3.2.0		new
- PyYAML	5.4.1		upgrade
- requests	2.25.1		upgrade
- scikit-learn	0.24.2		upgrade
- scipy	1.7.0		upgrade
- seaborn	0.11.1		upgrade
- statsmodels	0.12.2		upgrade
- tensorflow	2.5.0		upgrade
- Theano	1.0.5		
R	4.1.0		upgrade
- caret	6.0.88		upgrade
- knitr	1.33		upgrade
- rmarkdown	2.9		upgrade
- tidymodels	0.1.3		upgrade
- rsample	0.1.0		upgrade
- parsnip	0.1.6		upgrade
- recipes	0.1.16		upgrade
- workflows	0.2.2		upgrade
- tune	0.1.5		upgrade
- yardstick	0.0.8		upgrade

Component	Version	Remark	Status
- broom	0.7.8		upgrade
- dials	0.0.9		
- tidyverse			
- ggplot2	3.3.5		upgrade
- dplyr	1.0.7		upgrade
- tidyr	1.1.3		upgrade
- readr	1.4.0		
- purrr	0.3.4		
- tibble	3.1.2		upgrade
- stringr	1.4.0		
- forcats	0.5.1		upgrade
RStudio	1.4.1717	for virtual machine only	upgrade
Ubuntu	20.04.2 LTS	focal	upgrade
Vim	8.2.3083		upgrade
wget	1.21.1		

## Version 1.0.4

Component	Version	Remark	Status
Anaconda	1.7.2	2020.07	
Docker Compose	1.27.4		
Docker Engine	19.03.13		
Eclipse	2020-09	for virtual machine only	
Git	2.29.2		upgrade
Gradle	6.7		
ImageMagick	7.0.10-37		upgrade
Java (openjdk)	15	build 15+36-1562	
Python	3.8.3		
- pip	20.2.4		
- alpha_vantage	2.2.0		
- fbprophet	0.7.1		new

Component	Version	Remark	Status
- fire	0.3.1		new
- Keras	2.4.3		upgrade
- matplotlib	3.3.3		new
- numpy	1.19.4		new
- pandas	1.1.4		new
- PyYAML	5.3.1		new
- requests	n/a		new
- scikit-learn	0.23.2		upgrade
- scipy	1.5.4		upgrade
- seaborn	0.11.0		
- statsmodels	0.12.1		upgrade
- tensorflow	2.3.1		
- Theano	1.0.5		upgrade
R	4.0.3		
- caret	6.0.86		
- knitr	1.30		
- rmarkdown	2.5		
- tidymodels	0.1.1		
- rsample	0.0.8		
- parsnip	0.1.4		
- recipes	0.1.15		upgrade
- workflows	0.2.1		
- tune	0.1.1		
- yardstick	0.0.7		
- broom	0.7.2		
- dials	0.0.9		
- tidyverse			
- ggplot2	3.3.2		
- dplyr	1.0.2		
- tidyr	1.1.2		

Component	Version	Remark	Status
- readr	1.4.0		
- purrr	0.3.4		
- tibble	3.0.4		
- stringr	1.4.0		
- forcats	0.5.0		
RStudio	1.3.1093	for virtual machine only	
Ubuntu	20.04.1 LTS	focal	
Vim	8.1		

### Version 1.0.3

Component	Version	Remark	Status
Anaconda	1.7.2	2020.07	
Docker Compose	1.27.4		
Docker Engine	19.03.13		
Eclipse	2020-09	for virtual machine only	
Git	2.29.0		
Gradle	6.7		
ImageMagick	7.0.10-34		
Java (openjdk)	15	build 15+36-1562	
Python	3.8.3		
- alpha_vantage	2.2.0		
- keras	2.4.0		new
- pip	20.2.4		
- scikit-learn	0.23		new
- scipy	1.5.3		new
- seaborn	0.11.0		new
- statsmodels	0.12.0		new
- tensorflow	2.3.1		new
- theano	1.0.0		new
R	4.0.3		

Component	Version	Remark	Status
- caret	6.0.86		
- knitr	1.30		
- rmarkdown	2.5		
- tidymodels	0.1.1		
- rsample	0.0.8		
- parsnip	0.1.4		
- recipes	0.1.14		
- workflows	0.2.1		
- tune	0.1.1		
- yardstick	0.0.7		
- broom	0.7.2		
- dials	0.0.9		
- tidyverse			
- ggplot2	3.3.2		
- dplyr	1.0.2		
- tidyr	1.1.2		
- readr	1.4.0		
- purrr	0.3.4		
- tibble	3.0.4		
- stringr	1.4.0		
- forcats	0.5.0		
RStudio	1.3.1093	for virtual machine only	
Ubuntu	20.04.1 LTS	focal	
Vim	8.1		

## Version 1.0.2

Component	Version	Remark	Status
Anaconda	1.7.2	2020.07	
Docker Compose	1.27.4		
Docker Engine	19.03.13		

Component	Version	Remark	Status
Eclipse	2020-09	for virtual machine only	
Git	2.29.0		
Gradle	6.7		
ImageMagick	7.0.10-34		
Java (openjdk)	15	build 15+36-1562	
Python	3.8.3		
- alpha_vantage	2.2.0		new
- matplotlib	3.3.2		new
- pandas	1.1.3		new
R	4.0.3		
- caret	6.0.86		new
- knitr	1.30		new
- rmarkdown	2.5		new
- tidymodels	0.1.1		new
- rsample	0.0.8		new
- parsnip	0.1.4		new
- recipes	0.1.14		new
- workflows	0.2.1		new
- tune	0.1.1		new
- yardstick	0.0.7		new
- broom	0.7.2		new
- dials	0.0.9		new
- tidyverse			new
- ggplot2	3.3.2		new
- dplyr	1.0.2		new
- tidyr	1.1.2		new
- readr	1.4.0		new
- purrr	0.3.4		new
- tibble	3.0.4		new
- stringr	1.4.0		new

Component	Version	Remark	Status
- forcats	0.5.0		new
RStudio	1.3.1093	for virtual machine only	
Ubuntu	20.04.1 LTS	focal	
Vim	8.1		

---

### Version 1.0.1

Component	Version	Remark	Status
Anaconda	1.7.2	2020.07	
Docker Compose	1.27.4		upgrade
Docker Engine	19.03.13		
Eclipse	2020-09	for virtual machine only	
Git	2.29.0		upgrade
Gradle	6.7		upgrade
ImageMagick	7.0.10-34		new
Java (openjdk)	15	build 15+36-1562	
Python	3.8.3		
R	4.0.3		
RStudio	1.3.1093	for virtual machine only	
Ubuntu	20.04.1 LTS	focal	upgrade
Vim	8.1		

---

### Version 1.0.0

Component	Version	Remark	Status
Anaconda	1.7.2	2020.07	
Docker	19.03.13		
Docker Compose	1.27.1		
Eclipse	2020-09	for virtual machine only	
Git	2.28.0		
Gradle	6.6.1		
Java (openjdk)	15	build 15+36-1562	



Component	Version	Remark	Status
Python	3.8.3		
R	4.0.3		
RStudio	1.3.1093	for virtual machine only	
Ubuntu	20.04	focal	
Vim	8.1		

## 2. Creating a new Machine Learning development container

### 2.1 Getting started

```
> REM Assumptions:
> REM   - the name of the Docker container should be: my_ml_dev
> REM   - the path the host repository is: //C/projects/my_repro
> REM   - the directory name for this repository inside the container should be:
my_repro_dir
> REM   - you want to use the latest version of the Konnexions development image
> docker run --name my_ml_dev \
    -v //C/projects/my_repro:/my_repro_dir \
    konnexionsgmbh/ml_dev:latest

> REM Stopping the container
> docker stop my_ml_dev

> REM Restarting the container
> docker start my_ml_dev

> REM Entering a running container
> docker exec -it my_ml_dev bash
```

### 2.2 Detailed Syntax

A new container can be created with the **docker run** command.

#### Syntax:

```
docker run -it
    [--name <container_name>] \
    konnexionsgmbh/ml_dev[:<version>]
    [<cmd>]
```

#### Parameters:

- **container\_name** - an optional container identification
- **directory\_repository** - an optional host repository directory - the default value is expecting the repository inside the container
- **version** - an optional version number of the image or the constant **latest**
- **cmd** - an optional command to be executed in the container, default is **bash** for running the **bash** shell

Detailed documentation for the command **docker run** can be found [here](#).

#### Examples:

1. Creating a new Docker container named **my\_ml\_dev** using a repository inside the Docker container:

```
docker run -it --name my_ml_dev konnexionsgmbh/ml_dev:latest
```

2. Creating a new Docker container named **my\_ml\_dev** using the host repository of a Windows directory **D:\projects\my\_repro**:

```
docker run -it --name my_ml_dev -v //D/projects/my_repro:/my_repro konnexionsgmbh/ml_dev:latest
```

3. Creating a new Docker container named **my\_ml\_dev** using the host repository of a Linux directory **/my\_repro**:

```
docker run -it --name my_ml_dev -v /my_repro:/my_repro konnexionsgmbh/ml_dev:latest
```

## 3. Working with an existing Machine Learning development container

### 3.1 Starting a stopped container

A previously stopped container can be started with the **docker start** command.

#### Syntax:

```
docker start <container_name>
```

#### Parameter:

- **container\_name** - the mandatory container identification, that is an UUID long identifier, an UUID short identifier or a previously given name

Detailed documentation for the command **docker start** can be found [here](#).

### 3.2 Entering a running container

A running container can be entered with the **docker exec** command.

#### Syntax:

```
docker exec -it <container_name> <cmd>
```

**Parameter:**

- **container\_name** - the mandatory container identification, that is an UUID long identifier, an UUID short identifier or a previously given name
- **cmd** - the command to be executed in the container, e.g. **bash** for running the **bash** shell

Detailed documentation for the command **docker exec** can be found [here](#).

## 4. Best practices

### 4.1 Use of a root repository directory on the host computer

If all relevant repositories are located within a common parent directory, then development work in all these repositories can be done within a single Konnexions development container.

**Example:**

In the following example we assume that the host directory is named **C:\Temp\my\_projects** and should be mapped to the **projects** directory in the container.

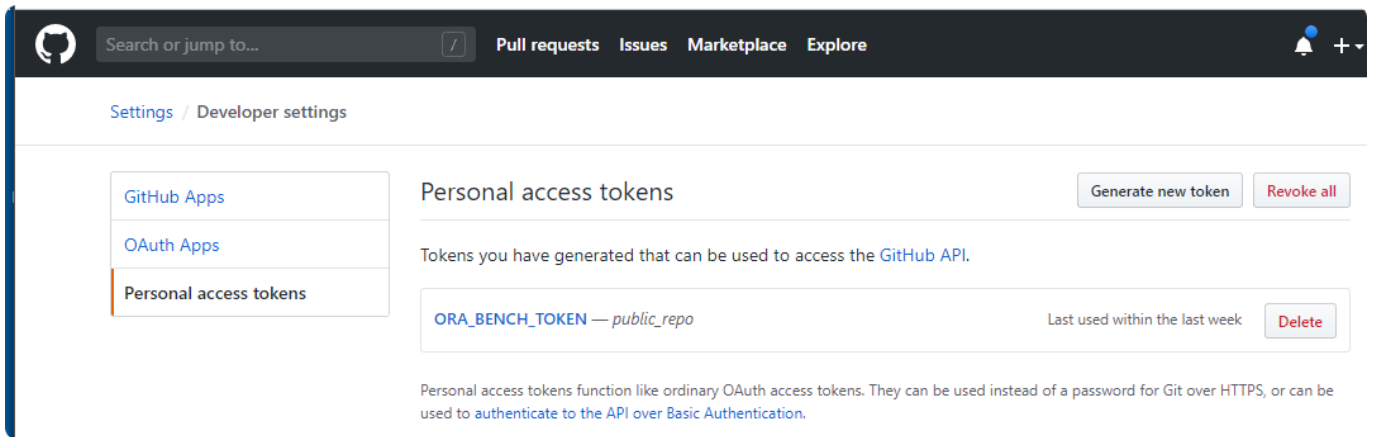
```
>C:\Temp\my_projects>docker run -it --name ml_dev -v  
//C/Temp/my_projects:/projects konnexionsgmbh/ml_dev:latest  
root@35b9310932f1:/# cd projects  
root@35b9310932f1:/projects# ls -ll  
total 0
```

### 4.2 Use of private GitHub repositories inside the container

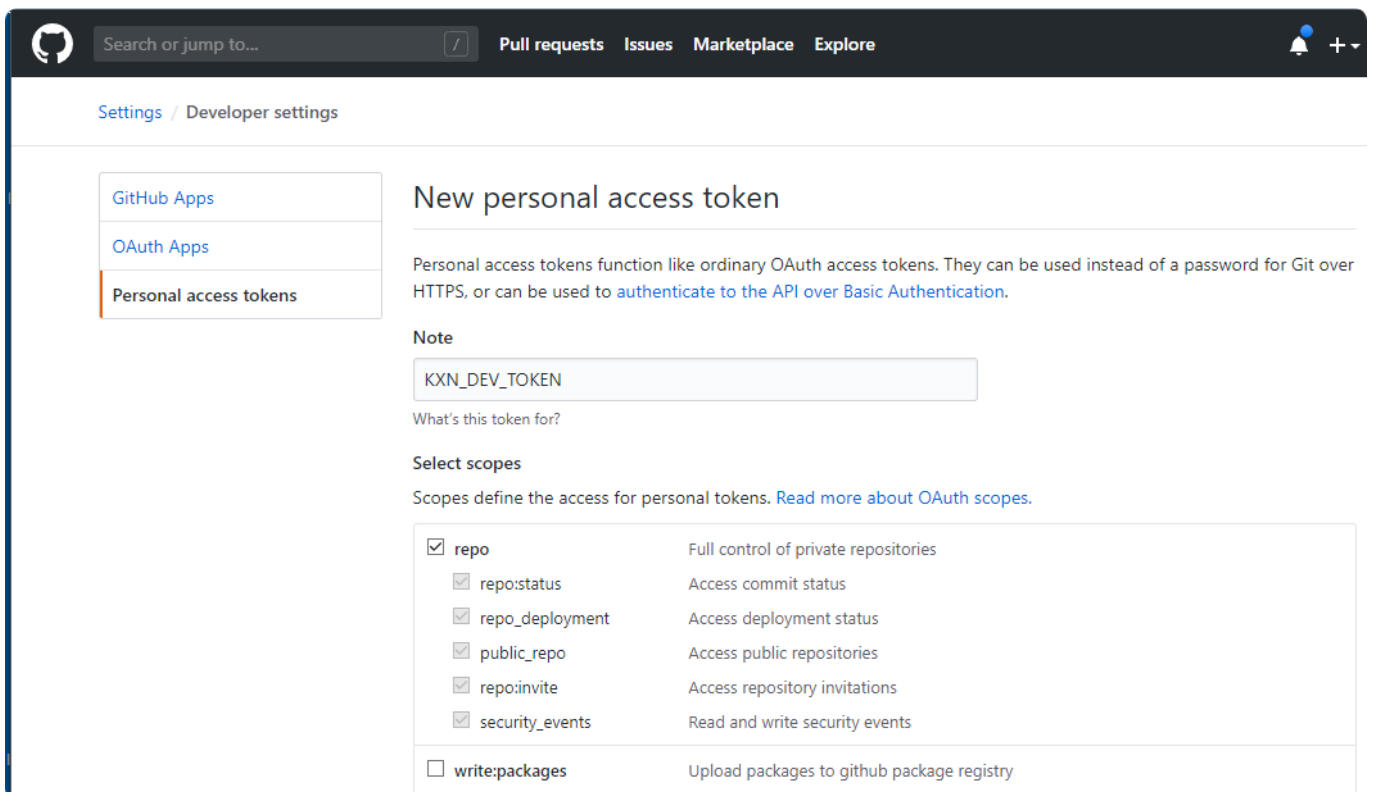
To access private repositories in GitHub, you must first create a new personal access token in GitHub and then add it to your git configuration inside the container.

#### 1. Create a new personal access token in GitHub

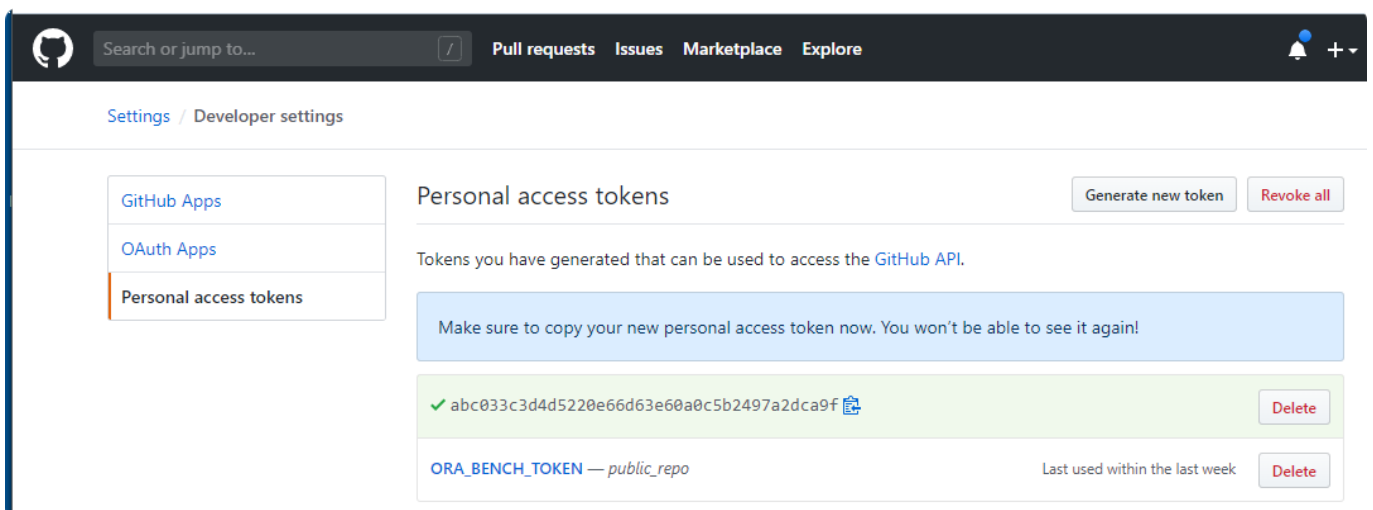
- With the following URL you can create the access token: <https://github.com/settings/tokens>



- Press the button **Generate new token**



- Name the new token, select the scopes and press the button **Generate token**



- Write down the secret code and keep it in a safe place

## 2. Setting up the Docker container on the host machine

In the following example we assume that the host directory is named `C:\Temp\my_projects` and should be mapped to the `projects` directory in the container.

```
C:\Temp\my_projects>docker run -it --name ml_dev -v //C/Temp/my_projects:/projects
konnexionsgmbh/ml_dev:latest
Unable to find image 'konnexionsgmbh/ml_dev:latest' locally
latest: Pulling from konnexionsgmbh/ml_dev
d51af753c3d3: Pull complete
...
a6bb30d1a5cf: Pull complete
Digest: sha256:5f6d6afc566ef9142d2d85b85dd331c0558eafaaf286179fd0ae787988c1b89b
Status: Downloaded newer image for konnexionsgmbh/ml_dev:latest
```

## 3. Initial configuration of git in the container

```
root@332206c300f1:/# export XDG_CONFIG_HOME=/projects
root@332206c300f1:/# mkdir -p $XDG_CONFIG_HOME/git/
root@332206c300f1:/# touch $XDG_CONFIG_HOME/git/config
root@332206c300f1:/# touch $XDG_CONFIG_HOME/git/credentials
root@332206c300f1:/# git config --file=$XDG_CONFIG_HOME/git/config
credential.helper 'store --file=/projects/git/credentials'
root@332206c300f1:/# git config --file=$XDG_CONFIG_HOME/git/config user.name "John
Doe"
root@332206c300f1:/# git config --file=$XDG_CONFIG_HOME/git/config user.email
"john.doe@company.com"
root@332206c300f1:/# git config --list --show-origin
file:/projects/git/config          credential.helper=store --
file=/projects/git/credentials
file:/projects/git/config          user.name=John Doe
file:/projects/git/config          user.email=john.doe@company.com
```

## 4. Verification of the settings

```
root@332206c300f1:/# cat /projects/git/config
[credential]
    helper = store --file=/projects/git/credentials
[user]
    name = John Doe
[user]
    email = john.doe@company.com
```

## 5. Clone a repository for the first time

When prompted provide your github user name and the new personal access token from (1).

```
root@332206c300f1:/# cd projects
root@332206c300f1:~# git clone https://github.com/KonnexionsGmbH/docker_images
Cloning into 'docker_images'...
Username for 'https://github.com': John Doe
Password for 'https://john.doe@company.com':
abc033c3d4d5220e66d63e60a0c5b2497a2dca9f
remote: Enumerating objects: 78, done.
remote: Counting objects: 100% (78/78), done.
remote: Compressing objects: 100% (49/49), done.
remote: Total 78 (delta 33), reused 68 (delta 23), pack-reused 0
Receiving objects: 100% (78/78), 167.83 KiB | 867.00 KiB/s, done.
Resolving deltas: 100% (33/33), done.
```

## 6. Verify if the clone completed with success

```
root@332206c300f1:~# cat /projects/git/credentials
https://John Doe:abc033c3d4d5220e66d63e60a0c5b2497a2dca9f@github.com
```

## 7. Verification after a restart of the Docker container

```
C:\Temp\my_projects>docker start ml_dev
ml_dev
C:\Temp\my_projects>docker exec -it ml_dev bash
root@332206c300f1:/# export XDG_CONFIG_HOME=/projects
root@332206c300f1:/# git config --list --show-origin
file:/projects/git/config credential.helper=store --
file=/projects/git/credentials
file:/projects/git/config user.name=John Doe
file:/projects/git/config user.email=john.doe@company.com
```

## 8. Verification after the removal of the Docker container

- Deleting the Docker container and image

```
C:\Temp\my_projects>docker stop ml_dev
ml_dev

C:\Temp\my_projects>docker rm ml_dev
ml_dev

C:\Temp\my_projects>docker images
REPOSITORY          TAG                 IMAGE ID            CREATED
SIZE
```

```

konnexionsgmbh/ml_dev      latest      51757b5e414e      6 hours ago
3.71GB

C:\Temp\my_projects>docker rmi 51757b5e414e
Untagged: konnexionsgmbh/ml_dev:latest
Untagged:
konnexionsgmbh/ml_dev@sha256:5f6d6afc566ef9142d2d85b85dd331c0558eafaaf286179fd0ae787988c1b89b
Deleted: sha256:51757b5e414e5333ace7b163484c06e4685c29312ad09d5d7d648c6936011a60
...
Deleted: sha256:7789f1a3d4e9258fbe5469a8d657deb6aba168d86967063e9b80ac3e1154333f

```

- Recreating the Docker container (and image)

```

C:\Temp\my_projects>docker run -it --name ml_dev -v //C/Temp/my_projects:/projects
konnexionsgmbh/ml_dev:latest
Unable to find image 'konnexionsgmbh/ml_dev:latest' locally
latest: Pulling from konnexionsgmbh/ml_dev
d51af753c3d3: Pull complete
...
a6bb30d1a5cf: Pull complete
Digest: sha256:5f6d6afc566ef9142d2d85b85dd331c0558eafaaf286179fd0ae787988c1b89b
Status: Downloaded newer image for konnexionsgmbh/ml_dev:latest
root@ad1f036bbc44:/# export XDG_CONFIG_HOME=/projects
root@ad1f036bbc44:/# git clone https://github.com/KonnexionsGmbH/docker_images
Cloning into 'docker_images'...
remote: Enumerating objects: 78, done.
remote: Counting objects: 100% (78/78), done.
remote: Compressing objects: 100% (49/49), done.
remote: Total 78 (delta 33), reused 68 (delta 23), pack-reused 0
Receiving objects: 100% (78/78), 167.83 KiB | 895.00 KiB/s, done.
Resolving deltas: 100% (33/33), done.

```

- If we use the same path - where `git/config` and `git/credentials` exist - as in Step 4, `git` access (clone/push/pull) doesn't ask for username/password anymore.